

Belief-State Geometry in Transformers: Reproduction and Unsupervised Recovery via Observable Operator Models

1 Introduction

Shai et al.[1] show that a transformer trained on sequences from a hidden data-generating process linearly represents, in its residual stream, the geometry of *belief states*: the Bayesian posterior distributions over the generator’s hidden states. The geometry is characterised by the *mixed-state presentation* (MSP): each history w induces a belief $b(w)$ over the hidden states, and the set of reachable beliefs forms a characteristic structure in the probability simplex over the hidden states (for classical data-generating processes)[2].

This work reproduces the supervised results and contributes an *unsupervised* estimator that recovers the belief-state geometry from activations and the network’s own next-token probabilities alone, without ground-truth belief labels or knowledge of the generating process, using observable operator models (OOMs). We present supervised and unsupervised belief state recovery on two processes; *Mess3*, a 3-state dense process with Sierpinski fractal geometry, and *RRXOR*, a 5-state synchronising process with 36 discrete belief states.

Both processes are trained using the paper’s architecture ($d_{\text{model}}=64$, 4 layers, 1 head, $d_{\text{mlp}}=256$, ReLU, LayerNorm, SGD) for 10^6 steps. Throughout, we use prefix deduplication: since causal attention ensures identical contexts produce identical activations, each unique prefix appears once (weighted by its total frequency) rather than at every position. We weight all regressions by prefix probability $P(w)$, following [2], to give priority to the model’s learned geometry in high-probability regions.

A limitation of the supervised approach is requiring knowledge of the data-generating process, which prevents us from applying this approach to large language models. If we could recover the belief geometry in an unsupervised way this would address this limitation. If transformers truly represent belief-state geometry linearly in their residual streams, then this geometry should be recoverable without labels or access to the generating process. Success on these controlled experiments would suggest that the approach could extend to large language models.

2 Results: Supervised vs. Unsupervised Recovery

2.1 The unsupervised estimator

The estimator recovers the belief subspace from activations and the network’s own softmax, using the *observability matrix* of an observable operator model (OOM) [4], but estimated in activation space. We refer to this estimator as the spectral-OOM. This differs from Baum-Welch (EM) [3], which fits a generative HMM to the token sequence rather than reading the representation out of the network.

Linearising the belief update. The network encodes belief affinely, $a(w) = b(w)\Psi + c$, where Ψ is the linear embedding and c is a constant offset. Because every regression below is fit with an intercept (equivalently, on $P(w)$ -centered activations), the offset c is absorbed and the operators are estimated on the linear part $b(w)\Psi$. The Bayesian update is projective (Eq. (1)),

$$b(wx) = \frac{b(w)T^x}{b(w)T^x\mathbf{1}}, \quad P(x|w) = b(w)T^x\mathbf{1}. \quad (1)$$

a homography: a linear map $b(w) \mapsto b(w)T^x$ followed by a normalisation that divides by a linear functional of $b(w)$. The ratio is non-linear in $b(w)$, and hence in $a(w)$, so no fixed matrix relates consecutive activations across all prefixes w .

However, observe that the *normalising denominator is exactly the next-token probability* $P(x|w) = b(w)T^x\mathbf{1}$, which is the network’s softmax output. Multiplying both sides by $P(x|w)$ removes the non-linearity, giving us

$$b(wx)P(x|w) = b(w)T^x, \quad (2)$$

whose right-hand side is now linear in $b(w)$. Right-multiplying (2) by Ψ rewrites it in activation coordinates. On the left, $P(x|w)$ is a scalar and $b(wx)\Psi = a(wx)$, giving $P(x|w)a(wx)$. On the right, inserting $\Psi\Psi^+$ (the identity on the belief row-space when Ψ has full row rank d , i.e. the encoding is injective) regroups $b(w)\Psi$ into $a(w)$:

$$b(w)T^x\Psi = b(w)(\Psi\Psi^+)T^x\Psi = (b(w)\Psi)(\Psi^+T^x\Psi) = a(w)A_x, \quad A_x := \Psi^+T^x\Psi.$$

Equating the two sides yields a single fixed operator A_x relating every parent activation $a(w)$ to its rescaled child $P(x|w)a(wx)$ (where child activation means the appended-token activation $a(wx)$)

$$a(w)A_x \approx P(x|w)a(wx). \quad (3)$$

The relation is approximate because Ψ is unknown, belief is only approximately encoded, and A_x are estimated from finite data. The map $A_x = \Psi^+T^x\Psi$ is simply the hidden-state operator T^x expressed in activation coordinates. Crucially, the rescaling by the network’s own $P(x|w)$ is what absorbs the prefix-dependent normaliser that made the update projective, leaving a relation that is *linear* in the sense that one matrix A_x holds for all w . Each A_x is thus estimable by a single ($P(w)$ -weighted) linear regression over transition pairs $(a(w), a(wx))$ rescaled by $P(x|w)$, no supervised targets are required.

Recovering the subspace. The operators A_x of Eq. (3) advance the state *within* activation space; we also need a *readout* from a state to what it predicts. Let C be the ($P(w)$ -weighted) ridge regression of activations $a(w)$ onto the softmax row $P(\cdot|w)$, the one-step observable. It is linear in belief because $P(x|w) = b(w)T^x\mathbf{1}$ (Eq. (1)) is a linear functional of $b(w)$, so the full softmax row is a linear image of $b(w)$ and hence of $a(w)$. We fit C as a linear map rather than reuse the network’s own activation-to-softmax map because the latter is non-linear (the softmax) and so cannot be composed into the matrix products below. C is the best linear readout to the next-token probability simplex, and the SVD acts on matrices. Both C and $\{A_x\}$ are fit by the same regression from one pass over the prefix tree, independently and in either order. Their only targets are the network’s softmax and (rescaled) child activations.

The pair $(C, \{A_x\})$ is the activation-space analogue of the output map and dynamics of a linear system. C says what is seen now, A_x says how the state evolves, so A_xC is what is seen one

step ahead, $A_x A_y C$ two steps ahead, and so on. Composing the A_x preserves linearity in $b(w)$, so every such multi-step observable is again linear in belief. Stacking them up to a horizon L gives the *observability matrix*

$$\mathcal{O} = [C \mid (A_x C)_x \mid (A_x A_y C)_{x,y} \mid \cdots \mid (A_{x_1} \cdots A_{x_L} C)_{x_1, \dots, x_L}], \quad (4)$$

indexed over all words up to length L . This is the (word-indexed) control-theoretic observability matrix of the estimated OOM; L is taken large enough that $\mathcal{O}$ reaches its full rank, i.e. enough future to make the hidden states distinguishable.

The column space of $\mathcal{O}$ is the image of belief space under the encoding, so it spans exactly the belief-carrying directions of activation space, the *belief subspace*, isolating it from the rest of the residual stream. Its dimension, the numerical rank of $\mathcal{O}$, is the model order. For an observable (minimal) process this equals the number of hidden states, while a weakly observable process exposes only its observable part and can rank below the true state count (see Section 2.4, RRXOR). An SVD $\mathcal{O} = U \Sigma V^\top$ then does two things: the singular-value spectrum locates the effective rank d (the elbow giving the model order), and the leading left singular vectors $B = U_{:,1:d}$ furnish an orthonormal basis for the belief subspace, which we use to project activations. Recovering a state subspace via the SVD of a matrix of observable statistics is the standard spectral-learning route [6, 5]. There they factor the past $\times$ future Hankel matrix, whereas here we build its observability factor directly from learned operators in activation space. The procedure uses only activations, softmax, and the prefix tree.

Recovering the subspace from the past: the controllability factor. The observability matrix (4) stacks *futures*: from a state, what is predicted ahead. The dual object stacks *pasts*: which states are *reachable* by driving the process forward from its start. Recall the forward (un-normalised) update of Eq. (2), $b(wx)P(x|w) = b(w)T^x$. Iterating from the stationary root belief b_0 along a word $w = x_1 \cdots x_k$ telescopes the scalars into the joint probability,

$$b(w)P(w) = b_0 T^{x_1} T^{x_2} \cdots T^{x_k}, \quad P(w) = \prod_i P(x_i \mid x_1 \cdots x_{i-1}), \quad (5)$$

since each step contributes one conditional factor and the denominators telescope. The right-hand side is linear in b_0 : the un-normalised reachable state after word w is just b_0 pushed through the product operator $T^w := T^{x_1} \cdots T^{x_k}$.

In activation coordinates. Right-multiply (5) by Ψ and insert $\Psi^+ \Psi$ between every adjacent pair of operators (the identity on the belief row-space when Ψ has full row rank d):

$$b_0 T^{x_1} \cdots T^{x_k} \Psi = (b_0 \Psi) (\Psi^+ T^{x_1} \Psi) \cdots (\Psi^+ T^{x_k} \Psi) = a_0 A_{x_1} \cdots A_{x_k},$$

with $a_0 := b_0 \Psi$ the root activation and $A_x = \Psi^+ T^x \Psi$ the same operators already estimated in Eq. (3). Because the left side equals $a(w)P(w)$, this gives the reachability relation

$$a_0 A_{x_1} \cdots A_{x_k} \approx P(w) a(w), \quad (6)$$

the exact dual of (3): the observability blocks compose operators *on the right of C* (futures read out forward), whereas reachability composes the *same* operators *on the left of a_0* (pasts driven forward from the root). No new operator fits are needed, the controllability factor reuses $\{A_x\}$ and the single root activation a_0 .

The controllability matrix. Stacking the reachable activations over all words up to horizon L gives

$$\mathcal{R} = \begin{bmatrix} a_0 \\ (a_0 A_x)_x \\ (a_0 A_x A_y)_{x,y} \\ \vdots \\ (a_0 A_{x_1} \cdots A_{x_L})_{x_1, \dots, x_L} \end{bmatrix}, \quad (7)$$

each row the (un-normalised) activation reached by its word. Its *row* space is the span of reachable states in activation space, the reachable belief subspace; for a process whose start state explores all of belief space this coincides with the column space of $\mathcal{O}$, both equal the d -dimensional belief subspace, which is the activation-space statement of the factorisation $H = \mathcal{R} \mathcal{O}^\top$ at rank d .

Estimation. Unlike the observability factor, $\mathcal{R}$ requires no per-word regression: each row is a *deterministic* product of the already-fitted $\{A_x\}$ applied to a_0 . The only estimated inputs are a_0 (the $P(w)$ -weighted mean root activation, or the empirical activation at the start token) and the operators themselves. In practice the rows should be $P(w)$ -*weighted* exactly as the observability construction is, because (6) produces the *un-normalised* $P(w) a(w)$, so high-probability prefixes already carry larger norm and dominate the SVD, matching the weighting of $\mathcal{O}$ and of the supervised decode.

2.2 Metric-correct recovery of the belief subspace

The estimator of Section 2.1 makes three implicit choices that are harmless when belief is *strongly observable* but that systematically lose the weak belief modes when it is not. First, it ranks directions by the ordinary (Euclidean) singular values of the observability matrix $\mathcal{O}$, which conflate “large in activation norm” with “important for belief”. Second, it estimates the operators $\{A_x\}$ by independent per-symbol regressions and only then composes them, discarding the fact that they share a common invariant subspace. Third, it reads the model order d off an elbow in that Euclidean spectrum, which need not exist. On Mess3 all three choices are benign—belief aligns with the next-token observable, the spectrum has a clean elbow, and the per-symbol fits already suffice. On RRXOR each one fails, and the failures compound: the belief-discriminating directions are precisely those that move the next token least (*next-token degeneracy*), so they carry small Euclidean singular values, no elbow marks them, and no single operator exposes them.

This section replaces each Euclidean choice with its metric-correct counterpart. The construction is unchanged in spirit—it still uses only activations, the model softmax, and the prefix tree—but it (a) handles the affine offset correctly, (b) weights output directions by the Fisher metric of the next-token simplex, (c) ranks belief directions by basis-independent Hankel singular values obtained by balancing the controllability and observability Gramians, (d) recovers the operators’ *common* invariant subspace rather than fitting them independently, and (e) selects the model order by spectral projection, which is robust to the non-diagonalisable operator algebra of processes like RRXOR.

The affine offset. The encoding is affine, $a(w) = b(w)\Psi + c$ (Section 2.1). The rescaled relation (3) is linear in $b(w)$ only after the constant c is removed, but c cannot be removed by a single global re-centring of activations: the rescaling by $P(x|w)$ in the child target makes the offset’s

contribution prefix-dependent,

$$P(x|w) (a(wx) - \bar{c}) \neq P(x|w) a(wx) - \bar{c}, \quad (8)$$

so subtracting a fixed $\bar{c}$ from the rescaled child corrupts the fit. The correct remedy is to fit *every* regression below (each A_x , the readout C , and the subspace operators) *with its own intercept*. Each fit then absorbs its own offset, the operator A_x is estimated on the linear part $b(w)\Psi$ alone, and the intercept is discarded for the spectral analysis—the belief subspace is spanned by the linear maps, not by the offsets. We write $\tilde{a}(w)$ for the activation with its regression intercept removed and use it implicitly in what follows.

A Fisher-metric readout. The readout C regresses activations onto the next-token row $P(\cdot|w)$. Scored in the ordinary Euclidean output geometry, this treats all coordinates of the simplex alike. But belief states are points in a probability simplex, and the meaningful output geometry is the Fisher–Rao metric, under which two next-token distributions are far apart when they differ on a *rare* token. RRXOR’s next-token-degenerate belief states are Euclidean-close in output space yet genuinely distinct; a Fisher-weighted readout separates exactly these.

Concretely, for a next-token distribution $p = P(\cdot|w)$ over V tokens the multinomial Fisher information is $g(p) = \text{diag}(1/p) - \mathbf{1}\mathbf{1}^\top$; we use the diagonal surrogate $\text{diag}(1/p)$, which decouples the output coordinates and so reduces to V independent weighted ridge regressions. Writing $\text{sw}_w = P(w)$ for the existing context weight, output column k of C is fit with per-row weights

$$\text{sw}_w \cdot \frac{1}{\max(P(x=k|w), p_{\min})}, \quad (9)$$

with a small floor $p_{\min}$ to control near-zero token probabilities. The map C is still $(D \times V)$ and enters every construction below unchanged; only its fit is reweighted.

The two Gramians. The observability matrix $\mathcal{O}$ (4) stacks futures and the controllability matrix $\mathcal{R}$ (7) stacks pasts. Their Euclidean SVDs rank directions by raw magnitude. The correct objects are the two *Gramians*, formed in the right inner products: the $P(w)$ metric on the domain of contexts and an output metric on the range of futures. The controllability Gramian is the $P(w)$ -weighted second moment of the reachable activations,

$$W_c = \sum_w P(w) \tilde{a}(w)^\top \tilde{a}(w) = \mathcal{R}^\top D_c \mathcal{R}, \quad D_c = \text{diag}(P(w)), \quad (10)$$

a $(D \times D)$ symmetric positive-semidefinite matrix measuring how strongly each activation direction is *reached* by the process. The observability Gramian weights the future block of $\mathcal{O}$ by a range metric M_o ,

$$W_o = \mathcal{O} M_o \mathcal{O}^\top, \quad (11)$$

measuring how strongly each direction is *seen* in the future. Taking $M_o = I$ recovers the Euclidean case; the informative choices weight the length- k observable block of $\mathcal{O}$ by λ^k with $0 < \lambda < 1$, damping the short, next-token-biased futures so that the longer composites that carry the weak modes are not drowned out. Both Gramians are $(D \times D)$ regardless of how many pasts or futures are enumerated, so this step is cheap.

Balanced truncation and Hankel singular values. Ranking belief directions now means balancing reachability against observability. Take positive-semidefinite square roots $W_c = L_c L_c^\top$ and $W_o = L_o L_o^\top$ and form the SVD of their cross matrix,

$$L_c^\top L_o = U \Sigma V^\top, \quad \Sigma = \text{diag}(\sigma_1 \geq \sigma_2 \geq \dots). \quad (12)$$

The σ_k are the *Hankel singular values* of the estimated realisation. They are *basis-independent*—invariant under any invertible change of activation coordinates, unlike the raw singular values of $\mathcal{O}$ —and each one weighs a mode by *both* how reachable and how observable it is. A mode that is reached (the process visits those belief states, so it registers in W_c) but only weakly seen through the next-token readout (small in W_o) still receives a non-negligible σ_k , instead of being discarded as it is by the $\mathcal{O}$ -SVD. The rank- d belief subspace is projected by the leading balanced directions,

$$B = L_c U \Sigma^{-1/2} \big|_{:,1:d}, \quad s(w) = \tilde{a}(w) B, \quad (13)$$

and the geometry is read out from $s(w)$ exactly as before. The Hankel spectrum $\{\sigma_k\}$ is the honest measure of how much belief structure is present: the gap between the modes that stand above the finite-sample noise floor and those that fall below it separates what is recoverable from any observational estimator from what is genuinely sub-floor, a distinction the Euclidean spectrum cannot make.

Joint recovery of the operators. The operators $\{A_x\}$ share a common invariant subspace—the belief subspace—because each is $A_x = \Psi^+ T^x \Psi$ with the *same* Ψ . Fitting them independently and composing them, as in Section 2.1, throws this away. For RRXOR the cost is acute: the input-1 operator is nilpotent ($T_1^5 = 0$) and so individually carries almost no belief information, while the input-0 operator is defective; only *jointly* do $\{T_0, T_1\}$ span the dynamics. We therefore recover the common invariant subspace directly, seeking an orthonormal $Q \in \mathbb{R}^{D \times d}$ that is simultaneously near-invariant under every operator,

$$Q^* = \arg \min_{Q^\top Q = I_d} \sum_x \| (I - Q Q^\top) A_x Q \|_F^2, \quad (14)$$

whose objective is zero exactly when $\text{span}(Q)$ is invariant under all A_x . This is a joint (approximate) Schur problem rather than a joint diagonalisation: the operators are non-normal and defective, so a common eigenbasis need not exist, but a common invariant subspace does. Crucially, (14) is *magnitude-agnostic*; it does not prefer the large-eigenvalue (strongly observable) directions, which is what a dominant-subspace iteration would return and exactly the wrong bias for weak belief modes. The residual value of (14) at Q^* is itself diagnostic: if it remains large, the *estimated* operators do not share a clean common invariant subspace—an honest statement about estimation noise and defectiveness, not a quantity to be tuned away.

Model order by spectral projection. For a strongly observable process the order d is the elbow in the (now Hankel) spectrum. For a defective process there may be no elbow: the powers A^k of a defective operator carry polynomial-times-geometric envelopes $k^j \lambda^k$, so the “modes” smear across singular scales and fixing d by hand from a known state count is not target-free. The principled alternative separates the state space by *generalised eigenspace* rather than by singular magnitude, via the Riesz spectral projector. For the marginal activation-space operator $\hat{A} = \sum_x A_x$ and a contour Γ_λ enclosing only the eigenvalue cluster at λ ,

$$P_\lambda = \frac{1}{2\pi i} \oint_{\Gamma_\lambda} (zI - \hat{A})^{-1} dz, \quad (15)$$

is the projector onto the invariant subspace of that cluster; it is robust to defective and repeated eigenvalues, where an eigenvector basis fails. The belief subspace is the span of the retained projectors and the order is its rank,

$$\mathcal{B} = \bigoplus_{\lambda \in \text{kept}} \text{range}(P_\lambda), \quad d = \dim \mathcal{B}. \quad (16)$$

A genuine spectral projector satisfies $P_\lambda^2 = P_\lambda$ with $\text{tr } P_\lambda$ equal to the algebraic multiplicity it encloses; these identities are the verification that the contour encloses the intended cluster. Which clusters are “belief-bearing” is itself a question—the nilpotent part of RRXOR contributes an eigenvalue at 0, so the near-zero cluster may carry real belief structure rather than noise—and we report d both with and without it, treating the difference as a finding.

Summary. The refined estimator is the same observable-operator construction read out in the correct geometry: intercepts absorb the affine offset (8); the readout is Fisher-weighted (9); directions are ranked by the basis-independent Hankel singular values (12) of the balanced Gramians (10)–(11) rather than by Euclidean magnitude; the operators are recovered through their common invariant subspace (14); and the order is fixed by spectral projection (15). On Mess3, where belief is strongly observable, every refinement reduces to the original construction and the recovery is unchanged. On RRXOR the refinements are exactly what is needed to stop discarding the weakly observable belief directions—though no change of geometry can manufacture signal that the Hankel spectrum shows to be genuinely below the noise floor, and distinguishing those two regimes is precisely what the basis-independent spectrum now lets us do.

2.3 Mess3




Figure 1: **Mess3: ground truth, supervised, and unsupervised.** Left: ground-truth belief geometry. Centre: supervised decode from full 64-D residual stream ($R^2 = 0.999$). Right: unsupervised spectral-OOM subspace ($R^2 = 0.998$). The estimator recovers the fractal nearly identically.

For the Mess3 process, unsupervised recovery is nearly perfect. The observability spectrum drops cleanly at $d = 3$, and the recovered 3-D subspace decodes belief at $R^2 = 0.998$, almost as good as supervised. The geometry emerges similarly during training (Figure 2) and survives identical controls: both cross-validate on held-out prefixes and collapse under belief-label (supervised) or softmax (unsupervised) shuffle (Figure 3). The 30/70 split is a slight change from the 20/80 split in the original because the original split provides too little data for unsupervised learning; its training unit is a transition (parent + all children), while supervised trains on single prefixes.




Figure 2: **Mess3 emergence over training.** Top: supervised decode at steps 100, 1000, 5000, 10^6 . Bottom: spectral-OOB subspace activations (affine-aligned for display). Both show identical emergence trajectory.




Figure 3: **Mess3 controls: cross-validation and shuffle.** Top: supervised. Bottom: unsupervised. Both recover the fractal on held-out prefixes (CV) and collapse dramatically under shuffle control.

2.4 RRXOR

For RRXOR the gap to supervised recovery is substantial. The unsupervised subspace recovers the diamond geometry but more diffusely (Figure 4). We report two complementary scores: per-position R^2 (full point cloud) and state-level R^2 (36 belief state centers-of-mass). The unsupervised subspace gives per-position $R^2 = 0.65$ (supervised 0.91) and state-level $R^2 = 0.82$ (supervised 1.00). The $d = 5$ dimension was fixed *a priori* from the known RRXOR state count, since the singular

spectrum showed no clean elbow for this process.




Figure 4: **RRXOR: ground truth, supervised, and unsupervised.** Left: ground-truth 36 belief states (PC1/PC3 plane, display only). Centre: supervised decode (per-position $R^2 = 0.91$). Right: recovered subspace (per-position $R^2 = 0.65$, state-level $R^2 = 0.82$). The gap reflects weak observability.

This gap is structural, not an artefact of estimator quality. The root cause is rank deficiency in the observability matrix: RRXOR’s belief directions produce only faint singular values in the predictive observables (the multi-step version of the 0.33 next-token limitation), making them difficult to recover unsupervised. We tested observability at horizons $L = 3$ through $L = 15$ with $P(w)$ -weighting on the full length-10 enumeration. The state-level geometry R^2 at $d = 5$ shows that longer horizons do not improve recovery, best performance is at $L = 3$ ($R^2 = 0.815$). This confirms that observability is maximized at the shortest tested horizon.

The layerwise analysis and distance preservation are recovered in the supervised case (Figure 5). Belief is spread across all layers (no single layer encodes it well), and is preserved in distances better than next-token. Notably, the unsupervised representation reverses this (Figure 6). Next-token distances ($R^2 \approx 0.68$) are preserved better than belief ($R^2 \approx 0.50$). A possible explanation for this might be that the spectral-OOM subspace is too biased towards the next-token observable C . Since RRXOR’s belief geometry is not strongly correlated with next-token structure, this method may not be well-poised to recover the true belief geometry.

3 Conclusion

The belief-state geometry can be recovered unsupervised from activations and softmax alone, nearly perfectly for Mess3 ($R^2 \approx 0.998$), but only partially for RRXOR (state-level $R^2 \approx 0.82$). The spectral-OOM anchors to next-token predictions (the softmax observable), but RRXOR’s belief states exhibit significant next-token degeneracy because many geometrically distinct states share similar next-token distributions ($R^2 = 0.31$ from the original paper). It is unclear whether this can be fully addressed, as there are no obvious alternative observables satisfying the required properties (linearity in belief, empirical accessibility). A longer context window would reduce finite-sample error but not the underlying degeneracy. Success on Mess3 demonstrates the method’s viability where belief geometry aligns with observable structure, and a richly-sampled quantum or post-quantum process is a natural next step. Even if the method never scales to large models, it remains useful across the many small models whose belief geometry we can now study to learn what such geometry reveals about what a model has learned.




Figure 5: **RRXOR diagnostics: supervised.** Left: per-activation belief-reconstruction MSE. Centre/right: belief and next-token distance preservation via hexbin.




Figure 6: **RRXOR diagnostics: unsupervised.** Same metrics as Figure 5, but from the recovered spectral-OOM subspace. Belief spread is recovered, but next-token distances dominate.

A Observable Operator Models (OOMs): Spectral-OOM Estimator Details

A.1 The spectral-OOM estimator on transformer activations

From Section 2.1, the network encodes belief affinely: $a(w) = b(w)\Psi + c$. Working in $P(w)$ -centered coordinates (the offset c is absorbed by the regression intercept), the projective update (1) is non-linear in a , but its normaliser is precisely the softmax $P(x|w) = b(w)T^x\mathbf{1}$. Multiplying both sides by the normaliser gives $b(wx)P(x|w) = b(w)T^x$; right-multiplying by Ψ linearises the update:

$$a(w)A_x \approx P(x|w)a(wx), \quad A_x = \Psi^+T^x\Psi. \quad (17)$$

Building the observability matrix. Let C be the ridge regression of activations onto the softmax (the one-step observable), and $\{G_x\}$ the estimated transition operators (ridge regressions of activations onto $P(x|w)$ -child-activations, i.e. the empirical version of A_x). All regressions are $P(w)$ -weighted to bias the estimator toward high-probability prefixes. Because each G_x is linear in belief and composition preserves linearity, every block below is linear in $b(w)$. Stacking multi-step observables to horizon 3 gives the observability matrix

$$\mathcal{O} = [C \mid (G_x C)_x \mid (G_x G_y C)_{x,y} \mid (G_x G_y G_z C)_{x,y,z}], \quad (18)$$

indexed over all words up to length 3. This is the (word-indexed) control-theoretic observability matrix of the estimated OOM: with a single-symbol alphabet it reduces to the classical Kalman form $[C \mid CA \mid CA^2 \mid \dots]$, and the token-indexed family $\{G_x\}$ is the natural OOM/PSR generalisation. An SVD $\mathcal{O} = U\Sigma V^\top$ yields the orthonormal basis $B = U_{:,1:d}$, with the numerical rank d giving the model order.

What is and isn't unsupervised. Subspace choice, operator fit, and one-step R^2 are fully target-free (activations + softmax only). Order selection is target-free where the spectrum has a clean elbow (Mess3, $d = 3$); for the weakly observable RRXOR there is no elbow, and $d = 5$ is set a priori from the known state count (Section 2.3). The final placement onto the simplex and numerical comparisons use ground-truth beliefs, but only for display and scoring, not for fitting.

B Operator Rollouts

A stronger test of the subspace is the generative *operator rollout*. The activation-space operators A_x of Eq. (3) were used only to build the observability matrix and locate the subspace; they are not iterated. The rollout instead asks whether a self-contained dynamical system fitted *inside* the recovered subspace can regenerate the entire geometry from the root alone.

Three steps produce it. (i) *Operator-consistency refit*: the recovered basis B is good for decoding but is not exactly closed under the rescaled-transition relation, so operators fitted in it drift when iterated; we refine B by alternating least squares (ALS) into an operator-consistent basis B_{op} , alternating between fitting operators in the current basis and rotating the basis toward one those operators map into itself.

B.1 Results

Mess3: the rollout reproduces the full fractal at $R^2 \approx 1.00$ from both supervised and unsupervised subspaces (Figures 7 and 8). Mess3's true operators are diagonalisable strict contractions. The fractal is their attractor, so errors contract away and the rollout self-corrects




Figure 7: **Mess3 supervised subspace.** Ground truth | supervised decode | operator rollout ($R^2 = 1.00$).

RRXOR: the rollout collapses, failing to reach the outer apexes from both supervised and unsupervised subspaces (Figures 9 and 10). This is intrinsic to RRXOR's operator algebra: both of the transition operators T_x are non-diagonalisable; the input-1 operator is nilpotent ($T_1^5 = 0$), the input-0 operator is defective with a repeated eigenvalue, so the belief states are transient points the




Figure 8: **Mess3 spectral-OOM subspace.** Ground truth | recovered subspace (raw) | operator rollout via $\hat{A}_x$ ($R^2 = 1.00$).

dynamics passes through, not a contractive attractor it settles onto and self-corrects toward. The rollout operators $\hat{A}_x$, fitted to approximate this dynamics, reproduce the same behaviour. With no attractor to contract onto, iteration accumulates error and the extreme apex states are never reached.




Figure 9: **RRXOR supervised subspace.** Ground truth | supervised decode (COM-geom 0.95) | operator rollout (COM-geom 0.63; extreme states not reached).




Figure 10: **RRXOR spectral-OOM subspace.** Ground truth | recovered subspace (raw, state-level $R^2 = 0.66$) | operator rollout via $\hat{A}_x$ (state-level $R^2 = 0.74$; extreme states central).

References

- [1] A. S. Shai, S. E. Marzen, L. Teixeira, A. G. Oldenziel, P. M. Riechers. *Transformers represent belief state geometry in their residual stream*. NeurIPS 2024; arXiv:2405.15943.
- [2] P. M. Riechers, T. J. Elliott, A. S. Shai. *Neural networks leverage nominally quantum and post-quantum representations*. arXiv preprint arXiv:2507.07432, 2025.
- [3] L. E. Baum, T. Petrie, G. Soules, N. Weiss. *A maximization technique occurring in the statistical analysis of probabilistic functions of Markov chains*. Annals of Mathematical Statistics 41(1):164–171, 1970.
- [4] H. Jaeger. *Observable operator models for discrete stochastic time series*. Neural Computation 12(6):1371–1398, 2000.
- [5] M. Thon, H. Jaeger. *Links between multiplicity automata, observable operator models and predictive state representations — a unified learning framework*. JMLR 16:103–147, 2015.
- [6] D. Hsu, S. M. Kakade, T. Zhang. *A spectral algorithm for learning hidden Markov models*. J. Computer and System Sciences 78(5):1460–1480, 2012.
- [7] M. L. Littman, R. S. Sutton, S. Singh. *Predictive representations of state*. NeurIPS 2001.
- [8] J. P. Crutchfield, K. Young. *Inferring statistical complexity*. Physical Review Letters 63(2):105–108, 1989.
- [9] C. R. Shalizi, J. P. Crutchfield. *Computational mechanics: Pattern and prediction, structure and simplicity*. J. Statistical Physics 104:817–879, 2001.